

HousePriceChallenge.pdf

March 10, 2017

1 House Price Challenge

Kaggle Munich Meetup 2017-03-09
Woodmark Consulting AG

- Wolfgang Moser wolfgang.moser@woodmark.de
- Frank Ditsche frank.ditsche@woodmark.de
- Simon Heeg simon.heeg@woodmark.de

2 Preparation

2.1 Libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from scipy.stats import skew
from sklearn import linear_model
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import cross_val_score
from IPython.display import display
```

2.2 Load input data

```
In [2]: train = pd.read_csv("input/train.csv")
test = pd.read_csv("input/test.csv")
```



2.3 Check data dimensions

```
In [3]: print("Number of features in training set: {}".format(train.shape[1]))
        print("Number of training data entries: {}".format(train.shape[0]))
        print("Number of test data entries: {}".format(test.shape[0]))
```

Number of features in training set: 81

Number of training data entries: 1460

Number of test data entries: 1459

2.4 Where is the ID / the training prices?

```
In [4]: print("First column in both sets is: {}".format(train.columns[0]))
        print("Last column in training set is: {}".format(train.columns[-1]))
```

First column in both sets is: Id

Last column in training set is: SalePrice

2.5 What do the features look like?

```
In [5]: train.iloc[:, :10].head(5)
```

```
Out[5]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	\
0	1	60	RL	65.0	8450	Pave	NaN	Reg	
1	2	20	RL	80.0	9600	Pave	NaN	Reg	
2	3	60	RL	68.0	11250	Pave	NaN	IR1	
3	4	70	RL	60.0	9550	Pave	NaN	IR1	
4	5	60	RL	84.0	14260	Pave	NaN	IR1	

	LandContour	Utilities
0	Lvl	AllPub
1	Lvl	AllPub
2	Lvl	AllPub
3	Lvl	AllPub
4	Lvl	AllPub

```
In [6]: fig, axs = plt.subplots(1,2,figsize=(17,7))
```

```
train['LotArea'].plot.density(ax=axs[0])
train['LotArea'].plot.density(ax=axs[1])
```

```
print("Max. lot area: {} ftš".format(train['LotArea'].max()))
print("Mean lot area: {:.2f} ftš".format(train['LotArea'].mean()))
```

Max. lot area: 215245 ftš


Mean lot area: 10516.83 ftš

```
In [7]: axs[0].set_xlim([0,train['LotArea'].max()])
        axs[0].set_xlabel("ft^2")
        axs[0].set_ylabel("density")

        axs[1].set_xlim([30000,train['LotArea'].max()]);
        axs[1].set_ylim([0.,0.00000075])
        axs[1].set_xlabel("ft^2"); axs[1].set_ylabel("")
```

Out[7]: <matplotlib.text.Text at 0xb9a5550>

```
In [8]: plt.show()
```



2.6 Make index column the actual data frame index


```
In [9]: train.set_index('Id', inplace=True)
        test.set_index('Id', inplace=True)
```

2.7 Split off the price column from the training data

```
In [10]: train_price = train["SalePrice"]
         train.drop("SalePrice", axis=1, inplace=True)
```

2.8 Check value ranges of (numerical) features

```
In [11]: plt.rcParams['figure.figsize'] = (16, 6)
         train.boxplot(showfliers=False, rot=90)
         plt.show()
```



2.9 Deal with numerical features first

```
In [12]: # extract locations of numerical features
num_feat = (train.dtypes != "object").as_matrix()
print("Number of numerical features: {}".format(np.sum(num_feat)))

train_num = train.iloc[:, num_feat]
test_num = test.iloc[:, num_feat]
```

Number of numerical features: 36

2.9.1 Check positivity and fill missing values with column means

```
In [13]: print("All numerical values in training set positive? {}".format(not (train_num < 0).any().any()))
# fill missing values in training set with column means
train_num = train_num.fillna(train_num.mean())
train.iloc[:, num_feat] = train_num


print("All numerical values in test set positive? {}".format(not (test_num < 0).any().any()))
# fill missing values in test set with TRAINING column means
test_num = test_num.fillna(train_num.mean())
test.iloc[:, num_feat] = test_num
```

All numerical values in training set positive? True

All numerical values in test set positive? True

2.10 Check skewness (third central moment)

```
In [14]: ax = train_price.plot.density()
ax.set_xlim([0,train_price.max()])
ax.set_xlabel("price")
ax.set_ylabel("density")
plt.show()
```



```
In [15]: print("Numerical feature columns:")
display(train.columns[num_feat])
print("Skewness of numerical training features:")
display(skew(train_num))
```

Numerical feature columns:

```
Index([u'MSSubClass', u'LotFrontage', u'LotArea', u'OverallQual',
       u'OverallCond', u'YearBuilt', u'YearRemodAdd', u'MasVnrArea',
       u'BsmtFinSF1', u'BsmtFinSF2', u'BsmtUnfSF', u'TotalBsmtSF', u'1stFlrSF',
       u'2ndFlrSF', u'LowQualFinSF', u'GrLivArea', u'BsmtFullBath',
       u'BsmtHalfBath', u'FullBath', u'HalfBath', u'BedroomAbvGr',
       u'KitchenAbvGr', u'TotRmsAbvGrd', u'Fireplaces', u'GarageYrBlt',
       u'GarageCars', u'GarageArea', u'WoodDeckSF', u'OpenPorchSF',
       u'EnclosedPorch', u'3SsnPorch', u'ScreenPorch', u'PoolArea', u'MiscVal',
       u'MoSold', u'YrSold'],
      dtype='object')
```

Skewness of numerical training features:

```
array([ 1.40621011,  2.38249918, 12.19514213,  0.21672098,
        0.69235521, -0.61283072, -0.5030445 ,  2.67366126,
```

```

1.6837709 ,    4.25088802,    0.9193227 ,    1.52268809,
1.37534174,    0.81219427,    9.00208042,    1.36515595,
0.59545404,    4.09918567,    0.03652398,    0.67520283,
0.21157244,    4.48378409,    0.67564577,    0.64889763,
-0.66748815,   -0.3421969 ,    0.17979594,    1.5397917 ,
2.36191193,    3.08669647,   10.29375236,    4.11797738,
14.81313466,   24.45163962,    0.21183506,    0.09616958])

```

2.11 Log(1+p) transform of skewed features

```

In [16]: skewed = (np.absolute(skew(train_num)) > 1)
         train_num.iloc[:, skewed] = np.log1p(train_num.iloc[:, skewed])
         test_num.iloc[:, skewed] = np.log1p(test_num.iloc[:, skewed])
         train_price = np.log1p(train_price)

         print("Skeweness of numerical training features after transformation:")
         display(skew(train_num))

```

Skeweness of numerical training features after transformation:

```

array([ 2.48741218e-01, -8.91060059e-01, -1.37263272e-01,
        2.16720977e-01,  6.92355214e-01, -6.12830724e-01,
       -5.03044497e-01,  4.80625684e-01, -6.17774284e-01,
        2.52110019e+00,  9.19322702e-01, -5.14937258e+00,
        8.00317572e-02,  8.12194273e-01,  7.45264962e+00,
       -6.13394321e-03,  5.95454038e-01,  3.92902155e+00,
        3.65239844e-02,  6.75202835e-01,  2.11572442e-01,
        3.86543714e+00,  6.75645767e-01,  6.48897631e-01,
       -6.67488146e-01, -3.42196895e-01,  1.79795942e-01,
        1.53378802e-01, -2.33732497e-02,  2.11010418e+00,
        7.72702571e+00,  3.14717122e+00,  1.43483416e+01,
        5.16538998e+00,  2.11835060e-01,  9.61695796e-02])

```

2.12 Normalize numerical features

```

In [17]: scaler = StandardScaler().fit(train_num)
         train_num = scaler.transform(train_num)
         test_num = scaler.transform(test_num)

```

2.12.1 Apply the transformed values to the orginial sets

```

In [18]: train.iloc[:, num_feat] = train_num
         test.iloc[:, num_feat] = test_num

```

2.13 Transform categorical to numerical features

```
In [19]: train_test = pd.concat([train, test])
        train_test = pd.get_dummies(train_test)
        train_test.iloc[:,40:50].head(4)
```

```
Out[19]:
```

	MSZoning_RM	Street_Grvl	Street_Pave	Alley_Grvl	Alley_Pave	\
Id						
1	0	0	1	0	0	
2	0	0	1	0	0	
3	0	0	1	0	0	
4	0	0	1	0	0	

	LotShape_IR1	LotShape_IR2	LotShape_IR3	LotShape_Reg	LandContour_Bnk
Id					
1	0	0	0	1	0
2	0	0	0	1	0
3	1	0	0	0	0
4	1	0	0	0	0

2.14 Split sets again

```
In [20]: train = train_test.iloc[:train.shape[0], :]
        test = train_test.iloc[train.shape[0]:, :]
```

3 Finally we can start with machine learning

3.1 First try: Linear Regression

```
In [21]: lin_reg = linear_model.LinearRegression()
        scores = cross_val_score(lin_reg, train, train_price,
                                cv=5, scoring='neg_mean_squared_error')
        print("Mean of 5 CV sqrt MSE: {:.4f}".format(np.sqrt(-scores.mean())))
```

Mean of 5 CV sqrt MSE: 1392066367.1519

3.2 Second try: Ridge Regression

```
In [22]: ridge = linear_model.Ridge(alpha=10.)
        scores = cross_val_score(ridge, train, train_price,
                                cv=5, scoring='neg_mean_squared_error')
        print("Mean of 5 CV sqrt MSE: {:.4f}".format(np.sqrt(-scores.mean())))
```

Mean of 5 CV sqrt MSE: 0.1279

Ridge Regression seems to work better!

3.3 Fit estimator to training data

```
In [23]: ridge.fit(train, train_price)
```

```
Out[23]: Ridge(alpha=10.0, copy_X=True, fit_intercept=True, max_iter=None,  
              normalize=False, random_state=None, solver='auto', tol=0.001)
```

3.3.1 Apply estimator to test set and retransform predicted prices

```
In [24]: preds = ridge.predict(test)
```

```
In [25]: preds_price = np.expm1(preds)
```

3.3.2 Prepare format for submission and save as CSV

```
In [26]: test_results = pd.DataFrame({'SalePrice': preds_price,  
                                     'Id': test.index})  
        test_results.set_index('Id', inplace=True)  
  
        test_results.to_csv("output/test_results.csv")
```